

Recreating NeRF

This is a PyTorch implementation based on the paper: <https://arxiv.org/abs/2003.08934>. The code takes mostly after the official tiny nerf implementation: https://colab.research.google.com/github/bmild/nerf/blob/master/tiny_nerf.ipynb

```
In [ ]: !pip install numpy tqdm ipywidgets matplotlib
!pip3 install torch torchvision torchaudio
!pip install imageio
!pip install python-ffmpeg
!pip install imageio-ffmpeg
```

Imports + Data

```
In [1]: import torch
import torch.nn as nn
import matplotlib.pyplot as plt
import torch.nn.init as init
import time
import numpy as np
import torch.nn.functional as F
import os
from tqdm import tqdm
from ipywidgets import interactive, widgets

#Example data http://cseweb.ucsd.edu/~viscomp/projects/LF/papers/ECCV20/nerf/tiny_nerf_data.npz
```

Encoding + Model

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

```
In [14]: def encoding(x, L=10):
    rets = [x]
    for i in range(L):
        for fn in [torch.sin, torch.cos]: # Use torch functions
            rets.append(fn(2. ** i * x))
    return torch.cat(rets, dim=-1) # Concatenate along the last dimension

class NeRF(nn.Module):

    def __init__(self, pos_enc_dim = 63, view_enc_dim = 27, hidden = 256) -> None:
        super().__init__()

        self.linear1 = nn.Sequential(nn.Linear(pos_enc_dim, hidden), nn.ReLU())

        self.pre_skip_linear = nn.Sequential()
        for _ in range(4):
            self.pre_skip_linear.append(nn.Linear(hidden, hidden))
            self.pre_skip_linear.append(nn.ReLU())

        self.linear_skip = nn.Sequential(nn.Linear(hidden+pos_enc_dim, hidden), nn.ReLU())

        self.post_skip_linear = nn.Sequential()
        for _ in range(2):
            self.post_skip_linear.append(nn.Linear(hidden, hidden))
            self.post_skip_linear.append(nn.ReLU())

        self.density_layer = nn.Sequential(nn.Linear(hidden, 1), nn.ReLU())

        self.linear2 = nn.Linear(hidden, hidden)

        self.color_linear1 = nn.Sequential(nn.Linear(hidden+view_enc_dim, hidden//2), nn.ReLU())
        self.color_linear2 = nn.Sequential(nn.Linear(hidden//2, 3), nn.Sigmoid())

        self.relu = nn.ReLU()
        self.sigmoid = nn.Sigmoid()

    def forward(self, input):

        # Extract pos and view dirs
        positions = input[..., :3]
        view_dirs = input[..., 3:]

        # Encode
        pos_enc = encoding(positions, L=10)
        view_enc = encoding(view_dirs, L=4)

        x = self.linear1(pos_enc)
        x = self.pre_skip_linear(x)

        # Skip connection
        x = torch.cat([x, pos_enc], dim=-1)
        x = self.linear_skip(x)
```

```

x = self.post_skip_linear(x)

# Density prediction
sigma = self.density_layer(x)

x = self.linear2(x)

# Incorporate view encoding
x = torch.cat([x, view_enc], dim=-1)
x = self.color_linear1(x)

# Color Prediction
rgb = self.color_linear2(x)

return torch.cat([sigma, rgb], dim=-1)

```

Get Ray + Render

```

In [15]: def get_rays(H, W, focal, c2w):
        """
        Generate rays for a given camera configuration.

        Args:
            H: Image height.
            W: Image width.
            focal: Focal length.
            c2w: Camera-to-world transformation matrix (4x4).

        Returns:
            rays_o: Ray origins (H*W, 3).
            rays_d: Ray directions (H*W, 3).
        """
        device = c2w.device # Get the device of c2w

        # Convert focal to float32 before moving to device
        if isinstance(focal, np.ndarray):
            focal = torch.tensor(focal.item(), dtype=torch.float32, device=device)
        elif not isinstance(focal, torch.Tensor):
            focal = torch.tensor(focal, dtype=torch.float32, device=device)
        elif focal.device != device:
            focal = focal.to(device)

        i, j = torch.meshgrid(
            torch.arange(W, dtype=torch.float32, device=device),
            torch.arange(H, dtype=torch.float32, device=device),
            indexing='xy'
        )

        # Process in chunks if needed for extremely high resolutions
        dirs = torch.stack(
            [(i - W * 0.5) / focal, -(j - H * 0.5) / focal, -torch.ones_like(i, device=device)], -1
        )

        # Matrix multiply in smaller batches if needed
        rays_d = torch.sum(dirs[...], None, 1) * c2w[:3, :3], -1
        rays_d = rays_d.reshape(-1, 3)
        rays_o = c2w[:3, -1].expand(rays_d.shape)

        return rays_o, rays_d

def render_rays(network_fn, rays_o, rays_d, near, far, N_samples, device, rand=False, embed_fn=None, chunk=512):
    def batchify(fn, chunk):
        return lambda inputs: torch.cat([fn(inputs[i:i+chunk]) for i in range(0, inputs.shape[0], chunk)], 0)

    # Sampling
    z_vals = torch.linspace(near, far, steps=N_samples, device=device)

    if rand:
        # Only add noise to each z_val when rand is True
        noise = torch.rand(z_vals.shape[:-1] + (N_samples,), device=device) * (far - near) / N_samples
        z_vals = z_vals + noise

    # More efficient way to create points
    rays_o_shaped = rays_o.unsqueeze(1) # [batch, 1, 3]
    rays_d_shaped = rays_d.unsqueeze(1) # [batch, 1, 3]
    z_vals_shaped = z_vals.unsqueeze(-1) # [batch, N_samples, 1]

    # More memory-efficient ray point generation
    pts = rays_o_shaped + rays_d_shaped * z_vals_shaped # [batch, N_samples, 3]

    # Normalize view directions and expand more efficiently
    rays_d_norm = torch.nn.functional.normalize(rays_d, dim=-1)
    view_dirs = rays_d_norm.unsqueeze(1).expand_as(pts) # [batch, N_samples, 3]

    # Combine for network input
    input_pts = torch.cat([pts, view_dirs], dim=-1) # [batch, N_samples, 6]

    # Clear intermediate tensors to free memory
    del pts, view_dirs, rays_o_shaped, rays_d_shaped, z_vals_shaped

    # Use smaller chunk size for network evaluation

```

```

raw = batchify(network_fn, chunk)(input_pts)

# Release input_pts memory
del input_pts

# The rest of the function remains similar
sigma_a = raw[...,0] # Shape: [batch, N_samples]
rgb = raw[...,1:] # Shape: [batch, N_samples, 3]

# Improved volume rendering - compute dists more efficiently
dists = torch.cat([
    z_vals[..., 1:] - z_vals[..., :-1],
    torch.ones_like(z_vals[..., 1:]) * 1e10
], dim=-1)

alpha = 1. - torch.exp(-sigma_a * dists)
alpha = alpha.unsqueeze(-1) # [batch, N_samples, 1]

# Computing transmittance with a more memory-efficient approach
ones_shape = list(alpha.shape)
ones_shape[1] = 1
ones = torch.ones(ones_shape, device=device)

# Compute transmittance efficiently
T = torch.cumprod(
    torch.cat([ones, 1. - alpha + 1e-10], dim=1),
    dim=1
)[:, :-1] # [batch, N_samples, 1]

weights = alpha * T

# Compute final colors and depths
rgb_map = torch.sum(weights * rgb, dim=1)
depth_map = torch.sum(weights.squeeze(-1) * z_vals, dim=-1)
acc_map = torch.sum(weights.squeeze(-1), dim=-1)

return rgb_map, depth_map, acc_map

```

Train Loop

```

In [16]: def train(images, poses, H, W, focal, testpose, testing, device):
    print(f"Using device: {device}")
    model = NeRF().to(device)

    criterion = nn.MSELoss(reduction='mean')
    optimizer = torch.optim.Adam(model.parameters(), lr=5e-4)
    scheduler = torch.optim.lr_scheduler.ExponentialLR(optimizer, gamma=0.99)

    n_iter = 1000
    n_samples = 64
    i_plot = 50
    psnrs = []
    iternums = []
    t = time.time()

    # Reduce batch size or use smaller image dimensions for training
    ray_batch_size = 1024 # Smaller batch size for rays
    downscale_factor = 2 # Downscale images for training

    # Downscale dimensions for training
    h_train = H // downscale_factor
    w_train = W // downscale_factor

    # Convert data to tensors ONCE but keep on CPU
    images_tensor = torch.from_numpy(images).float()
    poses_tensor = torch.from_numpy(poses).float()

    for i in range(n_iter):
        # Select a random image for training
        img_i = np.random.randint(images.shape[0])

        # Get target image and pose
        target = images_tensor[img_i] # Keep on CPU initially
        pose = poses_tensor[img_i].to(device) # Move to device only when needed

        # Downscale for training
        if downscale_factor > 1:
            target_resized = F.interpolate(target.permute(2, 0, 1).unsqueeze(0),
                size=(h_train, w_train),
                mode='bilinear').squeeze(0).permute(1, 2, 0)
        else:
            target_resized = target

        # Get rays for the downsampled image
        rays_o, rays_d = get_rays(h_train, w_train, focal / downscale_factor, pose)

        # Use only a subset of rays for training (random sampling)
        select_inds = np.random.choice(rays_o.shape[0], size=ray_batch_size, replace=False)
        rays_o = rays_o[select_inds].to(device)
        rays_d = rays_d[select_inds].to(device)
        target_s = target_resized.reshape(-1, 3)[select_inds].to(device)

```

```

optimizer.zero_grad()

# Use smaller chunk size for processing
chunk_size = 512 # Smaller chunk size reduces memory usage
rgb, depth, acc = render_rays(model, rays_o, rays_d, near=2., far=6.,
                              N_samples=n_samples, device=device,
                              rand=True, chunk=chunk_size)

# Compute loss on the ray batch
loss = criterion(rgb, target_s)

loss.backward()
optimizer.step()

# Clean up to free memory
del rays_o, rays_d, rgb, depth, acc, target_s
torch.cuda.empty_cache() if torch.cuda.is_available() else None

if i % i_plot == 0:
    print(f'Iteration: {i}, Loss: {loss.item():.6f}, Time: {(time.time() - t) / i_plot:.2f} secs per iter')
    t = time.time()

# Evaluate on a subset of the test image to save memory
test_ray_batch_size = 4096 # Larger batch for test, but still limited
# In the training loop where we evaluate the test image:
with torch.no_grad():
    # Make sure testpose is on the correct device
    testpose_device = testpose.to(device)

    # Get full resolution rays for test image
    test_rays_o, test_rays_d = get_rays(H, W, focal, testpose_device)

    # Process test rays in smaller batches
    test_rgb_parts = []
    for j in range(0, test_rays_o.shape[0], test_ray_batch_size):
        end_idx = min(j + test_ray_batch_size, test_rays_o.shape[0])
        batch_o = test_rays_o[j:end_idx]
        batch_d = test_rays_d[j:end_idx]

        rgb_batch, _, _ = render_rays(model, batch_o, batch_d, near=2., far=6.,
                                      N_samples=n_samples, device=device, chunk=chunk_size)

        # Move results to CPU immediately
        test_rgb_parts.append(rgb_batch.cpu())

    # Clean up batch
    del batch_o, batch_d, rgb_batch
    if torch.cuda.is_available():
        torch.cuda.empty_cache()

    # Combine batches on CPU
    test_rgb = torch.cat(test_rgb_parts, dim=0)
    test_rgb = test_rgb.reshape(H, W, 3)

    # Make sure testing is on CPU for comparison
    testing_cpu = testing.cpu()

    # Compute PSNR on CPU to save memory
    test_loss = F.mse_loss(test_rgb, testing_cpu)
    psnr = -10. * torch.log10(test_loss)

    psnrs.append(psnr.item())
    iternums.append(i)

    # Plot
    plt.figure(figsize=(10,4))
    plt.subplot(121)
    plt.imshow(test_rgb.numpy())
    plt.title(f'Iteration: {i}')
    plt.subplot(122)
    plt.plot(iternums, psnrs)
    plt.title('PSNR')
    plt.show()

    # Clean up
    del test_rgb, test_rgb_parts, test_loss

# Step learning rate scheduler
if (i + 1) % 100 == 0:
    scheduler.step()

return model

```

Load Data

```

In [17]: # Load data
data = np.load('nerf_dataset/nerf_dataset.npz')
images = data['images'].astype(np.float32)
poses = data['poses'].astype(np.float32)
focal = data['focal'].astype(np.float32) # Convert focal to float32
H, W = images.shape[1:3]
print(images.shape, poses.shape, focal)

```

```

# Check if MPS is available
device = "cuda" if torch.cuda.is_available() else "cpu"
if torch.backends.mps.is_available():
    device = "mps"

# Split into training and test sets
# Use the last image as test image
test_idx = -1 # Use the last image as the test image
testing = images[test_idx]
testpose = poses[test_idx]

# Use all other images for training (excluding the test image)
if test_idx == -1:
    train_images = images[:-1]
    train_poses = poses[:-1]
else:
    # If using a specific index
    train_images = np.concatenate([images[:test_idx], images[test_idx+1:]], axis=0)
    train_poses = np.concatenate([poses[:test_idx], poses[test_idx+1:]], axis=0)

# Ensure images have 3 channels (RGB)
if train_images.shape[-1] > 3:
    print(f"Original image shape: {train_images.shape}, trimming to 3 channels")
    train_images = train_images[..., :3]
    testing = testing[..., :3]

# Display test image
plt.figure(figsize=(8, 8))
plt.imshow(testing)
plt.title("Test Image")
plt.axis('off')
plt.show()

# Convert tensors to the selected device with float32 dtype
testing = torch.from_numpy(testing).float().to(device)
testpose = torch.from_numpy(testpose).float().to(device)

# Now use train_images and train_poses in your training loop
print(f"Training set: {train_images.shape} images")
print(f"Test image: {testing.shape}")
(2, 1080, 878, 3) (2, 4, 4) 2501.6567

```

Test Image





Training set: (1, 1080, 878, 3) images
Test image: torch.Size([1080, 878, 3])



Train

```
In [18]: model = train(train_images, train_poses, H, W, focal, testpose, testing, device)
```



Using device: mps
Iteration: 0, Loss: 0.088453, Time: 0.01 secs per iter





Iteration: 50, Loss: 0.044452, Time: 0.69 secs per iter





Iteration: 100, Loss: 0.018682, Time: 0.68 secs per iter





Iteration: 150, Loss: 0.009853, Time: 0.68 secs per iter





Iteration: 200, Loss: 0.008927, Time: 0.68 secs per iter





Iteration: 250, Loss: 0.009086, Time: 0.68 secs per iter





Iteration: 300, Loss: 0.007254, Time: 0.68 secs per iter





Iteration: 350, Loss: 0.005293, Time: 0.68 secs per iter





Iteration: 400, Loss: 0.005389, Time: 0.68 secs per iter





Iteration: 450, Loss: 0.006313, Time: 0.69 secs per iter





Iteration: 500, Loss: 0.005410, Time: 0.68 secs per iter





Iteration: 550, Loss: 0.005165, Time: 0.68 secs per iter





Iteration: 600, Loss: 0.005095, Time: 0.68 secs per iter





Iteration: 650, Loss: 0.005213, Time: 0.68 secs per iter





Iteration: 700, Loss: 0.005444, Time: 0.68 secs per iter





Iteration: 750, Loss: 0.004502, Time: 0.68 secs per iter





Iteration: 800, Loss: 0.003979, Time: 0.68 secs per iter





Iteration: 850, Loss: 0.004116, Time: 0.68 secs per iter



Iteration: 900, Loss: 0.002715, Time: 0.68 secs per iter



Iteration: 950, Loss: 0.003464, Time: 0.68 secs per iter



Render Video

```
In [27]: # Transformation matrices in PyTorch
trans_t = Lambda t: torch.tensor([
    [1, 0, 0, 0],
    [0, 1, 0, 0],
    [0, 0, 1, t],
    [0, 0, 0, 1]
], dtype=torch.float32, device=device)

rot_phi = Lambda phi: torch.tensor([
    [1, 0, 0, 0],
    [0, torch.cos(phi), -torch.sin(phi), 0],
    [0, torch.sin(phi), torch.cos(phi), 0],
    [0, 0, 0, 1]
], dtype=torch.float32, device=device)

rot_theta = Lambda th: torch.tensor([
    [torch.cos(th), 0, -torch.sin(th), 0],
    [0, 1, 0, 0],
    [torch.sin(th), 0, torch.cos(th), 0],
    [0, 0, 0, 1]
], dtype=torch.float32, device=device)

# Pose function with spherical coordinates
def pose_spherical(theta, phi, radius):
    c2w = trans_t(radius)
    c2w = torch.matmul(rot_phi(torch.tensor([phi / 180. * np.pi], dtype=torch.float32, device=device)), c2w)
    c2w = torch.matmul(rot_theta(torch.tensor([theta / 180. * np.pi], dtype=torch.float32, device=device)), c2w)
    c2w = torch.tensor([[-1, 0, 0, 0], [0, 0, 1, 0], [0, 1, 0, 0], [0, 0, 0, 1]],
        dtype=torch.float32, device=device) @ c2w
    return c2w

# Function for rendering based on user input
def f(**kwargs):
    c2w = pose_spherical(**kwargs)
    rays_o, rays_d = get_rays(H, W, focal, c2w[:3, :4]) # Get rays (this is a placeholder)
    c2w, rays_o, rays_d = map(Lambda t: t.to(device), (c2w, rays_o, rays_d))
    with torch.no_grad():
        rgb, depth, acc = render_rays(model, rays_o, rays_d, near=2., far=6., N_samples=64, device=device) # Render rays
    rgb = rgb.reshape(H, W, 3).cpu().detach()
    img = torch.clamp(rgb, 0, 1).numpy() # Clamp RGB values between 0 and 1 and convert to numpy

    plt.figure(2, figsize=(20, 6))
    plt.imshow(img)
    plt.show()

# Interactive slider setup for theta, phi, and radius
sldr = Lambda v, mi, ma: widgets.FloatSlider(
    value=v,
    min=mi,
    max=ma,
    step=.01,
)

names = [
    ['theta', [100., 0., 360]],
    ['phi', [-30., -90, 0]],
    ['radius', [4., 3., 5.]],
]

interactive_plot = interactive(f, **{s[0]: sldr(*s[1]) for s in names})
output = interactive_plot.children[-1]
output.layout.height = '350px'
interactive_plot
```

```
Out[27]: interactive(children=(FloatSlider(value=100.0, description='theta', max=360.0, step=0.01), FloatSlider(value=-...
```

```
In [24]: frames = []
for th in tqdm(np.linspace(0., 360., 120, endpoint=False)):
    c2w = pose_spherical(th, -30., 4.)
    rays_o, rays_d = get_rays(H, W, focal, c2w[:3, :4])
    c2w, rays_o, rays_d = map(Lambda t: t.to(device), (c2w, rays_o, rays_d))
    with torch.no_grad():
        rgb, depth, acc = render_rays(model, rays_o, rays_d, near=2., far=6., N_samples=64, device=device)
    rgb = rgb.reshape(H, W, 3)
    frames.append((255*np.clip(rgb.cpu().detach().numpy(), 0, 1)).astype(np.uint8))

import imageio
f = 'video.mp4'
imageio.mimwrite(f, frames, fps=30, quality=7)
```

```
0%|          | 0/120 [00:00<?, ?it/s]
```

RuntimeError Traceback (most recent call last)

Cell In[24], line 7

```
5 c2w, rays_o, rays_d = map(lambda t: t.to(device), (c2w, rays_o, rays_d))
6 with torch.no_grad():
----> 7     rgb, depth, acc = render_rays(model, rays_o, rays_d, near=2., far=6., N_samples=64, device=device)
8     rgb = rgb.reshape(H, W, 3)
9     frames.append((255*np.clip(rgb.cpu().detach().numpy(),0,1)).astype(np.uint8))
```

Cell In[15], line 61, in render_rays(network_fn, rays_o, rays_d, near, far, N_samples, device, rand, embed_fn, chunk)

```
58 z_vals_shaped = z_vals.unsqueeze(-1) # [batch, N_samples, 1]
60 # More memory-efficient ray point generation
----> 61 pts = rays_o_shaped + rays_d_shaped * z_vals_shaped # [batch, N_samples, 3]
63 # Normalize view directions and expand more efficiently
64 rays_d_norm = torch.nn.functional.normalize(rays_d, dim=-1)
```

RuntimeError: MPS backend out of memory (MPS allocated: 18.08 GB, other allocations: 1.09 MB, max allowed: 18.13 GB). Tried to allocate 69 4.51 MB on private pool. Use PYTORCH_MPS_HIGH_WATERMARK_RATIO=0.0 to disable upper limit for memory allocations (may cause system failure).

```
In [ ]: f = 'video.gif'
imageio.mimwrite(f, frames, fps=30)
```

```
In [ ]: from IPython.display import HTML
from base64 import b64encode
mp4 = open('video.mp4','rb').read()
data_url = "data:video/mp4;base64," + b64encode(mp4).decode()
HTML("""
<video width=400 controls autoplay loop>
  <source src="%s" type="video/mp4">
</video>
""") % data_url
```